

NodeRoom

AI that does banker work — and gets it accepted.

The work that gets rejected

- A banker spends 2-4 weeks after a conference sector-classifying hundreds of companies — or needs clean diligence the same day, because the targets fly out in 48 hours. *your note*
- Even when AI does that analysis well (~58/100 on the tasks), 0% of the best model's work is client-ready. Built with 502 bankers; only ~16% was even a useful starting point. 
- The misses aren't IQ — they're diligence: invented numbers, dead formulas, broken audit trails. I saw exactly this leading LLMSuite adoption at JPMorgan. *your note*

Sources: BankerToolBench — Lau et al., Handshake AI + McGill (arXiv:2604.11304); best model GPT-5.4: 0% client-ready, ~16% useful starting point, 502 bankers

NodeRoom: the finance war room

- You're right that AI will do junior-banker work. The unsolved part isn't intelligence — it's making the output sendable. *your note*
- A live finance war room — spreadsheet, deck, notebook, chat — each run by a dedicated, harnessed agent. Ask in plain language; it writes evidence-bearing cells, each tied to the exact source line. *your note*
- Every number tied to its exact source — cell, page, line — so the work is auditable. Designed to be accepted on the first pass. *your note*
- The moat compounds as models improve: the source-tied audit trail and banker-style handoff are workflow + data lock-in, not a capability gap. Labs, Copilot, and CRMs won't build it. *your note*

A market we already have ground truth in

\$16B

Generative AI in financial services by 2030, growing ~39%/yr (Grand View Research). We enter where it's already real: startup & middle-market banking diligence and prospecting.

Source: Grand View Research — Generative AI in Financial Services Market (to ~\$16B by 2030, CAGR 39.1%)

Proof of execution — over time

- Shipped cited-sources: hand it a claim and a PDF, it boxes the exact source line — or refuses if the quote isn't there. The thesis, as an installable skill. 
- Shipped the PowerPoint honesty skill that built this very deck — every claim sourced or flagged, no invented numbers. 
- Built our own BankerToolBench-faithful eval harness — Docker sandbox, the in-benchmark financial tools (SEC EDGAR/VDR), Gandalf verifier — running end-to-end on a scaffold smoke. No client-ready benchmark score yet; the real task runs are the next milestone. 
- Live in production: the harness pulls the right skill on demand (e.g. cited-sources) — how cheaper models get to client-ready. 

Sources: `.claude/skills/cited-sources` — shipped, MIT, packaged · `.claude/skills/powerpoint` — this deck · BankerToolBench loop ledger (internal) — records the scaffold-smoke run, not an official benchmark score · shipped to production — skill RAG release (commit 66010a87)

How we make money

- Land with seats for diligence and prospecting teams; expand to usage on agent runs as work moves into the room. *your note*
- Margin lever: our harness makes cheaper open models client-ready — DeepSeek V4 Pro near \$0.44 / \$0.87 per million tokens, a fraction of frontier pricing. 
- Pricing not set yet — land design-partner pilots first, then price on proven acceptance. *your note*

Sources: DeepSeek V4 Pro pricing \$0.435/\$0.87 per 1M tokens (Morph LLM, 2026)

Why this team

- Solo founder today — led LLMSuite adoption at JPMorgan; banking, data engineering, and agentic AI in one person. *your note*
- War-room design informed by an SVB-originated AI-notebook startup and the Meta PQX QA team. *your note*
- I've done the diligence and prospecting work we're automating — that's the ground truth. *your note*
- No co-founders or advisors yet — finding the right banking design-partner and field advisors is part of the ask. *your note*

The ask

- See the proof yourself: install cited-sources, give it a claim and a PDF, and watch it box the source — or refuse. *your note*
- Intros to banking design partners — MDs and analysts who'll pressure-test the room on live deals. *your note*
- Proof-first — not raising on a number yet: design partners and banker acceptance first, then raise against a real eval delta (client-ready off 0%). *your note*

To verify before presenting

Nothing outstanding — every claim is verified or owned by you. 